



AI-Powered Supplier Risk Intelligence

Business Development Plan

Product & Market Position • Go-to-Market Strategy • Pricing & Unit Economics
Product Roadmap & Scale-Up Plan • Financial Plan • Risk Analysis • Milestones

Prepared by the SupplyPulse founding team
August 2026 • Version 1.0

Internal planning document. Figures are tagged throughout as **MEASURED**, **TARGET**, or **ESTIMATE** — see the legend on the Table of Contents page.

Table of Contents

1. Executive Summary	3
2. Company Overview	4
3. Product & Technology — What We Have Now	5
4. Market Analysis	6
5. Competitive Landscape	7
6. Business Development & Go-to-Market Strategy	8
7. Pricing Strategy	9
8. Product Roadmap & Future Scale-Ups	10
9. Operations Plan	11
10. Financial Plan	12
11. Risk Analysis & Mitigation	13
12. Milestones & KPIs	14
13. Appendices	15

Legend used throughout this document: **MEASURED** — instrumented in code or drawn from actual cash/spend records, not estimated. **TARGET** — a planned change (e.g. repriced plans, scan deduplication) not yet shipped or live. **ESTIMATE** — a directional, bottom-up estimate (market sizing, CAC/LTV) — not third-party audited research. Wherever a projection assumes the founders go full-time or the company takes on outside capital, it is explicitly labeled **ACCELERATED SCENARIO** and is not part of this plan's committed baseline.

1. Executive Summary

SupplyPulse is an AI agent that watches a company's suppliers for early signs of trouble — bankruptcy filings, factory fires, product recalls, labour strikes, lawsuits, credit downgrades — and turns scattered public information into a single 0–100 health score, updated daily, with trend history and threshold alerts by email and Slack. It exists because mid-sized companies find out a supplier is in trouble only when the shipment doesn't arrive, at which point the only remaining options are expensive: air freight, spot-market buying, or stopping the line.

The product is a working MVP today, not a concept: authentication, supplier management, the six-step AI scoring pipeline, trajectory charts, tiered billing, and daily automated scans are all live and functioning end-to-end, representing roughly 85–90% of core planned functionality. The company is pre-launch and pre-revenue — no customer has been onboarded, piloted, or charged yet. It was built entirely by its three Sri Lanka-based co-founders, unfunded, for under \$1,000 in total cash spend **MEASURED**.

Instrumenting real per-scan cost has surfaced the single most important fact in this plan: at current live pricing, the business loses money on its best-fit customer. A Pro-tier customer monitoring 25 suppliers on daily scans generates 750 scans per month at roughly \$0.11 per scan **MEASURED** — about \$82 of cost against \$19.99 of monthly revenue, a customer that is unprofitable by a factor of roughly four. This is not a rounding error to grow past; it is architectural. The fix has two parts: **scan deduplication** (scoring a canonical supplier once per period and fanning the result out to every customer who watches it, since supplier risk is not customer-specific) and **repricing** to roughly \$49/month for Pro and \$249/month for Enterprise **TARGET**, which together are projected to reach 75–85% contribution margin.

This plan is deliberately built around two tracks, and the distinction between them matters throughout the document. The **bootstrapped baseline** is what the company is actually committing to: self-funded, reinvestment-driven growth, with the business becoming self-sustaining at roughly ten paying customers once dedup and repricing ship, since the founders are not currently drawing salaries. A more ambitious **ACCELERATED SCENARIO** appears at several points in this document — reaching roughly \$20,000 MRR and 450 customers within 24 months — but this scenario assumes the founders go full-time and/or the company takes on outside capital later. It is presented as a labeled possibility, not a commitment this plan makes. Sri Lanka's cost structure is the strategic advantage underneath both tracks: the company earns in USD and spends in LKR, so it can sustain price points and reach profitability at customer counts that a US- or Europe-based competitor with equivalent burn simply cannot match.

Bottom-up market sizing puts the total addressable market at roughly \$115M at current pricing, rising to about \$285M once repriced **ESTIMATE**, inside a broader \$3–4B supply-chain-risk software category dominated by six-figure enterprise contracts that structurally cannot serve a 25-supplier buyer. This plan's three-year serviceable-obtainable-market target is 1,500–3,000 paying customers, or roughly \$1.5–3M ARR **ESTIMATE** — a modest outcome by venture-scale standards, and a genuinely profitable one at this company's cost base.

2. Company Overview

2.1 Founding thesis

Companies that depend on suppliers find out about trouble the same way everyone else does — from the news, after it is too late to do anything but pay a premium. The warning signs are almost always public first, but they are scattered across thousands of sources and nobody is reading all of them on a buyer's behalf. Most companies' actual current practice is nothing more than an ad hoc Google Alert and a spreadsheet, which fires on irrelevant name collisions, has no score, no history, and no threshold — everyone stops reading it within a month. SupplyPulse's founding bet is that this gap — not the enterprise incumbents — is the real competitor, and that it can be closed profitably at a price point six-figure platforms cannot follow into.

2.2 Founders

Founder	Focus area	Background
Lavindu Perera	Backend, scheduled scanning, agentic pipeline	Undergraduate, University of Sri Jayewardenepura (Applied Sciences); full-stack development intern
Rusira Sandul	Frontend, prompt sanitization, AI search	Undergraduate, University of Westminster & University of Sri Jayewardenepura; ~2 years' industry experience as an Associate Software Engineer
Niduka Jayathunga	Deployment, billing, notifications, database design	Undergraduate, University of Sri Jayewardenepura (Applied Sciences)

All three are technical founders who have worked together on projects and hackathons since 2024; SupplyPulse is their most substantial joint effort to date. There is no commercial or sales founder on the team — this is named explicitly in Section 11 (Risk Analysis) as the single largest capability gap the company is working to close, rather than a weakness to obscure.

2.3 Company at a glance

Stage	Pre-launch, pre-revenue MVP
Legal structure	Not yet incorporated — deferred to avoid ongoing compliance costs while pre-revenue and self-funded. Plan: incorporate a Sri Lankan private limited company, with the three founders as equal shareholders (33.3% each) and standard four-year vesting with a one-year cliff, before scaling revenue.
Location	Colombo, Sri Lanka
Team size	3 co-founders, currently part-time, all technical
Cash position	~\$1,500 USD (~LKR 450,000) MEASURED , contributed by the founders; no debt or outside liabilities

Monthly burn (today)	~\$50–150/month MEASURED — AI/search API usage and hosting only; founders are unpaid, so there is no salary cost in this figure
Capital efficiency to date	A complete working system — AI scoring agent, automated daily monitoring, alerting, full web application, Stripe billing — built for under \$1,000 total cash spend MEASURED

3. Product & Technology — What We Have Now

3.1 The scoring pipeline

A single pipeline run executes per supplier, in six steps:




Figure 1 — the six-step scoring pipeline. Runs in under 60 seconds end-to-end.

Every score is source-linked — signals carry the article, filing, or report they came from, so the output is never a bare number without evidence behind it. The pipeline is a thin serverless application on rented infrastructure: a Next.js 16 web app plus roughly 15 Netlify serverless functions, with Supabase for authentication and Postgres, and a nightly cron job that rescans every active supplier automatically — there are no servers, no data centre, and no ops team to run.

3.2 Data model

Six core tables: `suppliers`, `supplier_scores`, `supplier_signals`, `alert_log`, `user_subscriptions`, `user_profiles`. Score history is append-only, which is what makes 30/60/90-day trajectory charts possible and gives every customer a signal history a competitor cannot backfill if they switch away.

3.3 Feature completeness

Feature	Status
Authentication & supplier CRUD	Live
6-step AI scoring pipeline	Live
Trajectory charts (30/60/90-day)	Live
Email + Slack threshold alerts	Live
Stripe billing, 3-tier enforcement	Live
Daily automated cron scanning	Live
Prompt-injection defence on scraped web content	Live
Dual search provider with automatic failover	Live
Per-user alert routing (currently one global inbox)	Not yet — Phase 1
Automated tests / CI	Not yet — Phase 1
Async job queue for scoring (currently synchronous)	Not yet — Phase 1
Versioned SQL migrations	Not yet — Phase 1

These gaps are stated candidly because they are already fully scoped, not because they were discovered late — see Section 8 for the phased fix plan.

4. Market Analysis

4.1 Market sizing

Sizing is built bottom-up from the buyer profile, not taken from a syndicated report. The target buyer is a manufacturer or distributor with 20–500 employees, \$5M–\$200M revenue, and 5–50 suppliers, several single-sourced — typically a Head of Supply Chain, Procurement Manager, COO, or (at smaller firms) the founder, for whom supplier monitoring is roughly 5% of one person's job, which is why it doesn't get done.

Segment	Estimated qualifying firms
US manufacturers, 20–500 employees	~60,000
US wholesale distributors	~40,000
UK + EU equivalent (manufacturing-heavy: Germany, Italy)	~150,000
Canada + ANZ	~20,000
Total qualifying firms	~270,000

TAM / SAM / SOM — current pricing vs. target repriced



Figure 2 — market funnel. All figures **ESTIMATE**, directional bottom-up sizing, not audited third-party research.

Cross-checked top-down: supply-chain risk software is roughly a \$3–4B category; a \$285M under-served SMB tier inside it is about 8%, which is consistent with the bottom-up figure.

4.2 The expansion thesis — how this becomes a much larger market

\$285M is a modest venture outcome, reached deliberately without inflating the near-term number. There is a credible path toward \$1–2B, via two separate levers rather than customer-count growth alone: (1) **pricing** — per-seat or per-supplier pricing that raises blended ACV toward \$3–5k for larger accounts, and (2) **vertical widening** — the same product, unchanged, applied beyond manufacturing to any business that depends on suppliers: healthcare, construction, retail — a further 3–5x with no product change. Neither is a claim; both are exactly what the next 18 months of the roadmap (Section 8) are designed to test.

4.3 Regulatory tailwind

Germany's Supply Chain Act (LkSG) is already in force; the EU's CSDDD is phasing in. Both push documented supplier due diligence onto companies well below enterprise size — buyers who currently have no affordable way to comply. This reaches Sri Lankan exporters directly too: apparel, tea, rubber, and food-processing firms that sell into Western brands are increasingly asked to evidence their own supply chains, giving the company a regionally relevant wedge alongside the Western market (see Section 6 for how this shapes geographic sequencing).

5. Competitive Landscape

5.1 Named competitors

Competitor	Segment	Where they beat us	Where we beat them
Resilinc / Interos / Everstream	Enterprise, \$100k–500k/yr	Depth: n-tier mapping, licensed proprietary data, human analyst validation, enterprise trust	Price, time-to-value — a supplier name and 60 seconds vs. weeks of onboarding and a dedicated analyst
Prewave	Vienna-based, EU-regulatory-focused	Multilingual coverage at scale, established LkSG/CSDDD credibility, real enterprise customers, funding	Self-serve — they sell top-down to large European manufacturers; a 60-person distributor cannot sign up and get a score today
Craft.co	Supplier intelligence / data platform, mid-market reach	Breadth of company data (financials, headcount, ownership)	Focus — they are a data platform you must interpret; we deliver one score, one trend, one alert that tells you when to act
Dun & Bradstreet / Moody's supplier modules	Credit data incumbents	Brand trust, credit data nobody else has	Coverage of small private suppliers with no credit file — usually the riskiest ones — and operational signals (fires, strikes, recalls) credit data misses entirely
The real competitor: doing nothing	Google Alerts + a spreadsheet	Free and frictionless	Actually working — alerts with no score, no history, no threshold, and no signal-to-noise filtering stop getting read within a month. In practice, the sales effort is spent competing with inertia, not with Resilinc.

5.2 Positioning — three pillars

Pillar	Claim	Why it holds
Price	~100× cheaper than any real alternative	Enterprise tools start at \$30k–\$250k/yr; SupplyPulse starts free, paid tiers from ~\$49/mo. Incumbents cannot descend — their cost base is data licensing and implementation staff.
Speed	Score in 60 seconds, not 6 weeks	A supplier name is the entire setup. No onboarding call, no analyst, no implementation contract.
Trajectory	Trend matters more than the score	A score of 82 → 74 → 61 over three weeks says "call them today." A competitor a customer switches to cannot backfill that history — it compounds into a switching cost the longer a customer stays.

5.3 Why incumbents can't follow down-market

This is not a temporary pricing gap that closes on its own. Incumbents' cost structure — data licensing contracts and implementation staff — makes serving a 25-supplier buyer structurally unprofitable for them. The core product idea that makes this viable is the mirror image of theirs: **incumbents own the risk data; SupplyPulse doesn't need to**. An AI agent reading the open web on demand removes the six-figure data pipeline and the six-week implementation entirely, and Sri Lanka's cost base compounds this — the company can profitably sustain a price point that a US or European competitor would lose money at.

5.4 What protects this position over time

Stated plainly rather than overclaimed: over a 2-year horizon, not much protects this technically. The pipeline is search plus two model calls; a competent team could rebuild the core in a couple of months. What protects the company over 2 years is position and workflow, not code — accumulated score history for each customer's suppliers (which cannot be backfilled by a switch), the alert routine embedded in a procurement lead's week, and being the first recognizable name in a segment nobody currently serves. Over 5 years, three things compound into a real moat: **longitudinal data** (score histories paired with actual outcomes let the product validate and tune predictive accuracy in a way a 2031 entrant starts at zero on), **cost density** (once scanning is deduplicated, the marginal cost of serving the hundredth customer watching a common supplier approaches zero), and **distribution/workflow depth** (alert routing, Slack integration, and eventually procurement-system integrations become the tool a team already uses — switching means losing history and rebuilding a routine).

6. Business Development & Go-to-Market Strategy

6.1 Acquisition channels

Product-led, self-serve acquisition: content and SEO around supplier-risk topics, founder-led outbound, and a free tier that converts to paid as a customer's supplier count outgrows 3. There is no outbound sales team at this stage — none of the three founders has sold B2B software before, which is named directly as the company's primary commercial gap (Section 11).

6.2 Design-partner motion

The immediate go-to-market motion is recruiting 5–10 design partners — mid-sized manufacturers and distributors, starting in Sri Lanka where the founders can meet buyers face to face and iterate weekly — given free Pro access in exchange for direct, structured feedback ahead of general launch. Twenty structured discovery conversations are the near-term priority; as of this plan, the company has spoken to zero customers to date, and closing that gap is treated as more urgent than any other activity.

6.3 Geographic sequencing — the go-to-market scale-up path




Figure 3 — geographic sequencing.

Sri Lanka comes first for validation, not revenue — design partners are reachable in person and the local apparel/export sector has this problem twice over (as a buyer of upstream suppliers, and as a supplier itself being asked to prove due diligence to Western brands). The United States is the primary revenue target: the largest concentration of mid-sized manufacturers and distributors, the highest willingness to pay for SaaS, English-language sources the AI agent already reads natively, and a self-serve motion with no geographic delivery constraint. The UK follows for the same language and buying behaviour, with a favourable time-zone overlap with Sri Lanka for support. Germany is deliberately delayed until the product can support German-language sources — entering under-resourced would waste the strongest regulatory tailwind (LkSG) available. India/South Asia is a live option, not yet committed, tested only if the US self-serve motion proves harder than expected.

6.4 Local ecosystem — awareness only, not solicited here

As the company grows, Sri Lankan startup-support programs worth being aware of include ICTA initiatives, local accelerators such as Hatch, and angel networks such as the Lankan Angel Network. These are noted here as options for the future, not something this document requests.

7. Pricing Strategy

7.1 Current pricing and why it doesn't cover cost

Tier	Current price	Limits
Starter	Free	3 suppliers, weekly scans, email alerts — an acquisition channel, not a revenue line
Pro	\$19.99/mo	25 suppliers, daily scans, email + Slack alerts
Enterprise	\$99.99/mo	Unlimited suppliers, 6-hourly scans, priority support

The core pricing problem **MEASURED**: cost per scan is approximately \$0.11 — five Tavily web searches at ~\$0.016 each (\$0.08), GPT-4o-mini signal extraction (~\$0.003), and GPT-4o score synthesis (~\$0.025). A Pro customer watching 25 suppliers on daily scans generates 750 scans/month, roughly **\$82 of cost against \$19.99 of revenue** — not marginally unprofitable, underwater by roughly 4×. No price fix alone solves this at current scan volumes; the architecture has to change first.

7.2 The two-part fix

1. **Scan deduplication** **TARGET** — a supplier's risk profile is not customer-specific. Scoring a canonical supplier once per period and fanning the result out to every customer watching it turns cost from scaling with customers to scaling with distinct suppliers, a 5–20× reduction at any real density.
2. **Repricing** **TARGET** — Pro to ~\$49/month, Enterprise to ~\$249/month, once dedup ships.

Tier	Current	Target	Target contribution margin
Starter	Free / 3 suppliers	Free / 3 suppliers (unchanged — acquisition funnel)	n/a
Pro	\$19.99/mo	~\$49/mo	75–85% TARGET
Enterprise	\$99.99/mo	~\$249/mo	

One additional cash-flow factor drives scan cadence design: free-tier users cost real money too — an estimated \$1.43/month per free account **MEASURED**; at 10,000 free accounts that is ~\$14,300/month. This is the specific reason free-tier scanning runs weekly rather than daily, and why dedup matters even before the paid tiers are fixed.

7.3 Positioning by audience

Audience	One-line pitch
Buyer (procurement/ops)	"Know if your supplier is in trouble before it hits your business — for less than a team lunch per month."
Partner / market context	"Supply chain risk management is a \$3–4B market. Nobody owns the self-serve, under-\$100/month tier. We do."
Channel partner	"Your clients already need this — add supplier risk monitoring to what you offer them in one integration."

8. Product Roadmap & Future Scale-Ups




Figure 4 — product roadmap. Phase 4 is demand-triggered, not scheduled — see 8.2.

8.1 Phases 1–3

Phase 1 fixes the known engineering gaps from Section 3 and ships scan deduplication — the precondition for positive margins, not an optimization. Phase 2 adds the trust layer that turns "AI said so" into an auditable claim: source links on every signal, a confidence score, sanctions-list checks (OFAC/EU), and "what changed since last scan" diffing. Phase 3 adds the features needed for a team to use the product together — multi-tenant/organization support, CSV bulk import, and digest emails.

8.2 Phase 4 — conditional, not committed

Phase 4 includes ERP/procurement-system integrations (SAP, Coupa, NetSuite). This appears to conflict with Section 8.3's exclusion list below, which also names ERP integrations — it does not. Phase 4 sits at the far edge of the roadmap specifically because it is **demand-triggered**: it moves up only when a specific paying customer's deal actually requires it, consistent with the exclusion list's intent. Until then it is correctly described as "later, and only if," not a scheduled deliverable.

8.3 What we are deliberately not building in the next 18 months

- **N-tier supply chain mapping** — the marquee enterprise feature, and it requires exactly the data-licensing pipeline the company refuses to build. Chasing it makes SupplyPulse a worse version of Resilinc.
- **ERP and procurement integrations** (SAP, Coupa, NetSuite) — every enterprise prospect asks; each is months of work and a nine-month sales cycle the company isn't built to run. The customer's system of record today is a spreadsheet.
- **Enterprise readiness** — SSO, SOC 2, RBAC, custom contracts — deferred until a specific paying customer requires it.
- **Buying proprietary data** — licensing a credit or news feed would reintroduce the fixed cost that prices incumbents out of this segment. Zero data licensing is the strategy, not a gap.
- **Training a proprietary ML model** — this is a workflow and data-history company, not an ML company. Frontier models get cheaper and better every quarter at no cost to the business.
- **Adjacent risk categories** — ESG scoring, compliance reporting, forced-labour audit trails. Real markets with genuine regulatory tailwind, but a different buyer and a different sale; taking them on before the core product retains customers would fragment a three-person team.

The trigger that would change this list: if inbound demand from 500+ employee companies becomes the majority of signups, ERP integrations and enterprise readiness move up the list. The intent is to be pulled up-market by revenue, not pushed there on a hunch.

8.4 Technical scale-up path — hyperscale architecture (long horizon, not a near-term commitment)

This subsection describes what changes if the company reaches roughly 100,000 users and \$1M+ MRR — a future that implies a funded or profitable team, not something to build prematurely. It is included here because the user asked for future scale-ups to be part of this plan, not because it is scheduled.

- **Dedicated worker fleets** (ECS Fargate / Railway / Kubernetes) pulling from a real queue (SQS / Kafka / Redis Streams), once serverless invocation costs or timeouts start to bite — likely somewhere in the 5,000–20,000 user range, well before hyperscale.
- **Database evolution** — read replicas separating dashboard reads from the write-heavy scan pipeline, and monthly partitioning on `supplier_scores` / `supplier_signals` as history grows into tens of millions of rows.
- **Shared signal-cache layer** — this is the same mechanism as the near-term scan-deduplication fix in Section 7, at a larger scale: the marginal cost of serving an additional customer watching an already-covered supplier approaches zero.
- **Batch AI APIs** for scheduled (non-interactive) scans — roughly a 50% cost reduction with no quality loss, since cron scans aren't latency-sensitive.
- **Full observability** — APM/tracing, per-customer cost dashboards — necessary at this volume to catch an unprofitable or abusive account before it costs real money.
- **SOC 2 Type II** — pursued only once an actual enterprise deal requires it, not speculatively.

9. Operations Plan

9.1 Team and structure

Three co-founders, currently part-time alongside degree studies and internships/employment, each owning a distinct area (backend/pipeline, frontend/AI search, deployment/billing/DB — see Section 2.2). All three are prepared to move to full-time as the business is able to support it. There is no non-founder hire yet; the first planned hire is growth/sales, not engineering, since that is the acknowledged capability gap.

9.2 Technology stack

Next.js 16, Supabase (Postgres + Auth), Stripe, OpenAI (GPT-4o / GPT-4o-mini), Tavily with Serper as automatic failover, Netlify (hosting, serverless functions, scheduled jobs), Resend and Slack webhooks for alert delivery — see Section 3 for the full architecture.

9.3 Operating costs — today vs. the accelerated scenario

	Today (bootstrapped baseline)	ACCELERATED SCENARIO
Team cost	\$0 — founders unpaid MEASURED	3 founders full-time: ~\$4,500–5,500/month, rising to ~\$7,000/month with a part-time commercial hire
Infrastructure / AI / search	\$50–150/month MEASURED	Same base, scaling with volume and free-tier usage
Total monthly burn	~\$50–150/month	~\$4,500–7,000/month

9.4 Payments and entity structure

Stripe handles all card data, keeping PCI scope out of the company's own systems. One unresolved operational question: revenue arrives in USD via Stripe while most costs are in LKR, and appropriate banking arrangements are needed to receive and hold foreign currency efficiently, along with confirmation of whether Stripe can settle directly to a Sri Lankan entity or whether an offshore structure is required. This is flagged candidly as open rather than assumed away — see Section 11.

10. Financial Plan

10.1 Unit economics

Cost component (per scan)	Amount MEASURED
5 Tavily web searches (~\$0.016 each)	\$0.08
GPT-4o-mini signal extraction	\$0.003
GPT-4o score synthesis	\$0.025
Total cost per scan	~\$0.11 (falling to ~\$0.078 at volume pricing)

	Today	Repriced TARGET
Annual value per Pro customer	\$240/yr	\$588/yr
Gross LTV (24-month retention assumption)	~\$480 ESTIMATE	~\$1,175 ESTIMATE
Blended CAC (self-serve/content-led)	\$150–350 ESTIMATE (assumes free-to-paid conversion of 3–5%)	

10.2 Two growth tracks

This is the section where the distinction between the two tracks matters most — neither table below should be read in isolation from its label.

Bootstrapped baseline — what this plan commits to. Self-funded, reinvestment-driven. Founders continue unpaid; monthly burn stays at \$50–150/month MEASURED. Once dedup and repricing ship, the business becomes self-sustaining at roughly **10 paying customers** — a low bar that does not depend on any external funding, and is reachable purely from the design-partner motion in Section 6.

ACCELERATED SCENARIO — contingent on the founders going full-time and/or the company taking on capital later. Not committed to by this plan.

Milestone	MRR	Paying customers	Team	Monthly burn
Month 6	~\$700	~15	3 founders	~\$4,500
Month 12	~\$4,500	~100	3 founders	~\$6,500
Month 18	~\$11,000	~250	3 founders + 1 growth hire	~\$7,500
Month 24	~\$20,000 (~\$240k ARR)	~450	4–5	~\$7,000+

Under this scenario, covering a ~\$7,000/month cost base at repriced levels requires roughly 175–200 paying customers (vs. 800+ a comparably-staffed US-based team would need), with operating cash-flow breakeven targeted around month 20–22. The single assumption most likely to be wrong is free-to-paid conversion (assumed 3–5%); at 1% these targets roughly halve, and the company's low burn means that adjustment is survivable rather than fatal.

10.3 Cost-side scale-up note

The hyperscale infrastructure described in Section 8.4 would only step-change the cost structure well beyond this plan's 24-month horizon — roughly \$1M+ MRR territory. It is not modeled into either track above; it is flagged here only as a forward-looking consideration.

10.4 Open financial questions

- Sri Lanka foreign-exchange and inward-remittance rules for USD SaaS revenue — banking arrangements and Central Bank compliance not yet finalized.
- Export-service tax treatment — software delivered to overseas customers should qualify for export-service treatment; to be confirmed with local counsel.

11. Risk Analysis & Mitigation

Risk	Category	Impact	Mitigation
No founder has sold B2B software before	Market	High	Design-partner motion generates real discovery-call reps before any paid acquisition spend; first non-founder hire is growth/sales, not engineering
Zero customer validation to date	Market	High	20 structured discovery conversations and 10 design partners are the immediate, sequenced first priority (Section 6.2)
Well-funded competitor (Prewave) could move down-market	Market	Medium	Channel conflict protects this position — an incumbent's own sales team and pricing structure make a \$49/mo product economically irrational for them to launch alongside a \$200k one
Known engineering gaps (single alert inbox, no CI/tests, synchronous scoring)	Technical	Medium	Fully scoped, first in Phase 1 (Section 8.1) — before any multi-tenant/public launch
AI-generated scores about real companies — defamation / accuracy exposure	Legal	Medium	Every signal is source-linked with evidence retained; outputs are kept factual and disclaimed as an indicator only, not financial/legal/investment advice; a supplier correction route is planned
GDPR — non-EU company processing EU-relevant data	Regulatory	Medium	Will require an EU representative and appropriate data-transfer safeguards; flagged as an area needing external guidance
Sri Lanka Personal Data Protection Act No. 9 of 2022	Regulatory	Low–Medium	Compliance alignment planned as the Act comes fully into force
EU AI Act transparency obligations	Regulatory	Low	Product use sits outside the high-risk category, but AI-generated-output disclosure is already standard product practice
Forex / inward remittance and export-tax treatment unresolved	Financial	Medium	Open question requiring local banking and tax guidance before revenue scales materially (Section 9.4, 10.4)
Current pricing is margin-negative if repricing/dedup slips	Financial	High	Both are first in Phase 1 — treated as a precondition to marketing spend, not a later optimization
Thin cash runway under the bootstrapped baseline (~\$1,500 on hand)	Financial	Medium	Runway is effectively open-ended at current burn since there is no fixed obligation — the constraint is pace, not survival

12. Milestones & KPIs

12.1 Near-term milestones — bootstrapped baseline

Milestone	Target
Discovery conversations	20 structured conversations with procurement/ops leads, first 3 months
Design partners	10 recruited, using the product weekly
First paying conversions	5 of the 10 design partners converted to paid within 6 months
Self-sustaining	Breakeven at ~10 paying customers post dedup/repricing

12.2 **ACCELERATED SCENARIO** milestones (labeled, contingent — restated from Section 10.2)

Milestone	Target
Gross margin fix	75%+ contribution margin via scan deduplication and repricing
Acquisition channel	CAC under \$350 with 3%+ free-to-paid conversion by month 12
Revenue	~\$20k MRR / ~\$240k ARR by month 24, annual retention above 85%
Operating breakeven	Month 20–22

12.3 KPIs tracked regardless of scenario

- Free-to-paid conversion rate
- Customer acquisition cost (CAC)
- Contribution margin %
- Monthly recurring revenue (MRR)
- Customer retention / churn (underlies the 24-month LTV assumption)
- Design-partner-to-paid conversion rate

13. Appendices

Appendix A — Regulatory & legal detail

Area	Status / obligation
Licensing to operate	None required — the company is not a credit rating agency or financial adviser; every score carries a disclaimer that it is an operational indicator only, not financial, legal, or investment advice
GDPR / UK GDPR	Data-processing agreements in place with sub-processors (Supabase, OpenAI, Stripe); an EU representative and transfer safeguards will likely be required as a non-EU company serving EU customers
Sri Lanka Personal Data Protection Act No. 9 of 2022	Applies as a locally established company; alignment planned as the Act comes fully into force
EU AI Act	Use sits outside the high-risk categories; transparency obligations (disclosing AI-generated output) are already met
Payments / PCI	Stripe handles all card data; PCI scope stays out of company systems
Foreign exchange / inward remittance	Open — banking arrangements and Central Bank compliance for USD revenue not yet finalized
Export-service tax treatment	Open — expected to qualify for export-service treatment, to be confirmed with local counsel
SOC 2 Type II	Not legally required; commercially deferred by design until a paying customer requires it

Appendix B — Competitor profiles (expanded)

Resilinc / Interos / Everstream — established enterprise supply-chain risk platforms; \$100k–500k/yr contracts; sell n-tier mapping, licensed proprietary data, and human analyst validation; weeks of onboarding with a dedicated analyst required; structurally cannot serve a 25-supplier buyer profitably.

Prewave — Vienna-based, well-funded, strong EU regulatory positioning and multilingual coverage; sells top-down to large European manufacturers via a sales-led motion with no self-serve path.

Craft.co — a broader supplier-intelligence data platform (financials, headcount, ownership) with mid-market reach; a data platform a buyer must interpret themselves, rather than a single score/trend/alert.

Dun & Bradstreet / Moody's — credit-data incumbents with strong brand trust; supplier risk is a module within a broader credit-data product; blind to small private suppliers with no credit file and to operational signals (fires, strikes, recalls) that credit data misses.

Appendix C — Glossary

Term	Definition
ACV	Annual Contract Value — the yearly revenue from one customer

ARR / MRR	Annual / Monthly Recurring Revenue
CAC	Customer Acquisition Cost
LTV	Customer Lifetime Value
TAM / SAM / SOM	Total / Serviceable Addressable / Serviceable Obtainable Market
Contribution margin	Revenue minus variable cost of serving a customer, before fixed costs
Scan deduplication ("dedup")	Scoring a canonical supplier once per period and reusing the result for every customer watching that supplier, instead of re-scanning per customer
Idempotency check	The pipeline step that prevents re-scanning a supplier that was already scanned in the current period
Worker fleet / queue	Dedicated background compute (vs. serverless functions) pulling jobs from a message queue — a scale-up pattern, not current architecture
SOC 2 Type II	An independent audit of a company's security controls over time, often required to close enterprise deals
GDPR / EU representative	EU data-protection law; non-EU companies processing EU personal data generally must appoint a local representative
LkSG / CSDDD	Germany's Supply Chain Due Diligence Act / the EU's Corporate Sustainability Due Diligence Directive — regulations requiring documented supplier due diligence

This document is an internal planning tool. Figures marked **ESTIMATE** are directional and have not been independently audited. Figures marked **ACCELERATED SCENARIO** assume a change in funding or staffing status this document does not commit to. SupplyPulse risk scores are indicators only, not financial, legal, or investment advice.