

Acornic Ventures · Startup Evaluation

SupplyPulse · MVP · IT · Logistics · 2026-07-29



About your startup

A few quick details so we know who you are.

What is your startup called?

SupplyPulse

Your name and role

Lavindu - co-founder

What stage are you at?

MVP

What sector or industry are you in?

IT

Which country or market are you focused on first?

Logistics

Today's date

2026-07-29

A. Your team

Tell us who is building this.

1. Who are the founders? For each person, give their name, role, and whether they work on this full-time or part-time.

Lavindu — Co-founder & Developer. Currently part-time; ready to go full-time on funding.

Rusira Sandul — Co-founder & Developer. Currently part-time; ready to go full-time on funding. Connect with me on rusirasandul.online .

Niduka — Co-founder & Developer. Currently part-time; ready to go full-time on funding.

All three are technical. We have no commercial founder

2. What experience does each founder have that is relevant to this business? (years, past companies, results)

We are early in our careers and this is our first venture, so we'd rather point to what we've built than to job titles.

What we've demonstrated: we designed and shipped a complete, working product with no funding — an AI agent pipeline (multi-source web search, signal extraction, score synthesis), a full Next.js application, Supabase auth and data, Stripe billing, scheduled scanning, and email/Slack alerting. It runs end-to-end today. That covers the full stack of what this business requires technically.

Specific relevant capability: production LLM engineering — cost-optimised model routing, prompt-injection defence for untrusted web content, multi-provider failover, and instrumented per-run unit economics. This is the hard technical core of the product and we built it ourselves.

Lavindu — An undergraduate from University Of Sri Jayewardenepura Applied Sciences. Experience in Software Engineering as currently a fullstack developing intern. Leads the backend, schedule scan, & agentic pipeline.

Rusira — [X years] in [area]. Leads the frontend, prompt sanitization, & AI search.

Niduka — An undergraduate from University Of Sri Jayewardenepura Applied Sciences. Leads the deployment, billings, notifications, & DB design.

What we lack is commercial experience.

3. Have any of the founders worked together before? On what?

Yes, we have worked together on projects and hackathons since 2024. SupplyPulse is our latest project together.

Building this product to a working state has already tested how we split work and make decisions: Lavindu owns the backend, schedule scan, & agentic pipeline, Rusira the frontend, prompt sanitization, & AI search Niduka owns the deployment, billings, notifications, & DB design.

4. Who are the founders and what are their shareholdings? What are the shareholdings of others, if any? Provide details of any ESOP in place.

Not yet formalised, the company is unincorporated. On incorporation we intend an equal three way split (33.3% each) among the founders. There are no other shareholders, advisers holding equity, or outstanding options.

5. Who owns the product or technology (the intellectual property) right now?

The codebase and product were built solely by the three founders, using no employer resources, and are currently held jointly by us as individuals. There are no third party contributors, contractors, or open source obligations beyond standard permissive licence dependencies.

On incorporation, all three founders will execute IP assignment agreements transferring full ownership of the code, product and brand to the company.

6. Is anything holding a founder back — a current job, non-compete, or visa issue?

All three founders currently work part time on SupplyPulse alongside degrees, and internships, which is the main constraint on our pace. All three are prepared to transition to full-time on funding.

No non-compete restrictions apply, our current work is unrelated to supply chain software.

No visa or right-to-work issues.

7. What key skills is the team missing, and how will you fill them?

We are three developers. The gaps are entirely on the commercial side, and we'd rather name them precisely than pretend otherwise.

Sales and go-to-market is the critical gap. None of us has sold B2B software. We can build a product that scores a supplier in 60 seconds; we don't yet know how to get a procurement manager at a 100 person manufacturer to take a call. How we fill it: mentorship and founder-led sales coaching through the Acornic programme in the first 3–6 months, with the aim that at least one founder can run discovery and close independently. Our first non-founder hire will be growth/sales, not engineering.

Domain expertise in procurement. We understand the problem from research, not from having lived it. How we fill it: recruiting design partners who become de facto advisers, and ideally an adviser with a procurement or supply-chain background — an introduction we'd value highly.

Marketing and distribution. No experience with SEO, content, or paid acquisition in B2B. How we fill it: structured experiments in the first six months with programme guidance, then hire once we know which channel works.

Finance and commercial operations. Pricing strategy, forecasting, and fundraising mechanics are new to us. How we fill it: mentorship, plus external accounting once incorporated.

What we're not missing: engineering, AI/LLM systems, product build, or infrastructure. We've shipped all of it without funding. The technical risk is largely retired; the remaining risk is commercial, and that's exactly the gap the programme fills.

B. The problem

What problem are you solving, and for whom?

8. What exact problem are you solving, and who has this problem?

The problem: companies find out a supplier is in trouble only when the shipment doesn't arrive.

The warning signs are almost always public first — bankruptcy filings, factory fires, product recalls, labour strikes, lawsuits, credit downgrades, a plant in a region hit by sanctions or a typhoon. That information exists in news, legal filings and financial coverage. It is simply scattered across thousands of sources, and nobody is reading it on your behalf for the 25 companies you depend on.

By the time you find out, the cheap options are gone. What remains is expensive: air freight, spot-market buying, or stopping the line.

Who has it: the procurement or operations lead at a small-to-mid-sized company that manufactures or distributes physical goods. Typically 20–500 employees, \$5M–\$200M revenue, 5–50 suppliers with several single-sourced. Titles vary — Head of Supply Chain, Procurement Manager, COO, sometimes the founder. Critically, there is no risk analyst: supplier monitoring is about 5% of one person's job, which is why it doesn't get done.

Best-fit segments: contract manufacturers and electronics assemblers, where one component supplier stops a line; food and CPG brands with co-packers and ingredient suppliers, where recall exposure is direct; distributors and e-commerce sellers with concentrated overseas sourcing.

This problem is on our doorstep too. Sri Lanka's export manufacturers — apparel, tea, rubber, food processing — have it twice over. They depend on upstream suppliers of fabric, dyes, chemicals and packaging, and they simultaneously sell to Western brands who are increasingly required to document supplier due diligence under EU regulation. That gives us a pool of design partners we can meet in person, and a live regional wedge alongside the Western market.

Who does not have this problem: large enterprises (already have the tooling and the analysts), pure software and services businesses (no physical supply chain), and anyone with two suppliers they speak to weekly.

9. How do people deal with this problem today, and why is that not good enough?

1. Nothing — reactive. The most common approach by far. You find out when the supplier calls, or when the PO slips. By then your only moves are expensive: air freight, spot-market buying, or a production stoppage.

2. Onboarding checks. A questionnaire, a credit report, maybe an audit — done once at contract signing, then never repeated until renewal. This treats risk as a fixed score when it is actually a moving one. A supplier that passed in January can be filing for bankruptcy in June.

3. Google Alerts and a spreadsheet. Free, so people try it. It fires on press releases and irrelevant name collisions, misses anything in the supplier's local-language press, and has no score, no history and no threshold. Everyone stops reading it within a month.

4. Enterprise platforms — Resilinc, Interos, Everstream, Prewave. These genuinely work. They also cost six figures a year, take weeks to onboard, and assume you have an analyst to run them. They are built for large buyers with thousands of suppliers.

The gap: options 1–3 are free but do not work. Option 4 works but is priced roughly 100x out of reach for a company with 25 suppliers. There is nothing in the middle, so mid-sized buyers default to hoping.

This is not a temporary gap. Incumbents cannot come down-market without breaking their own economics — their cost base is data licensing and implementation staff, so serving a small buyer is structurally unprofitable for them. The middle stays empty unless someone builds for it deliberately.

What is missing from all four options is the same thing: continuous monitoring and a trend. Even a good one-time assessment tells you a supplier is a B today. It does not tell you they were an A six weeks ago — and that movement is the signal that actually predicts a problem.

10. How many potential customers have you spoken to? What are the three most useful things you learned?

None yet — zero customer conversations to date.

We chose to build first, so that early conversations could open with a live risk score for a prospect's own supplier rather than a description of one. In hindsight that ordering is debatable, and correcting it is our immediate next step: 20 discovery conversations and 10 design partners within the first three months of the programme. We would begin with Sri Lankan export manufacturers, where we can meet buyers face to face and iterate weekly, and ideally within the Acorn Group ecosystem.

Since we have no customer learnings to report, the three most useful things we have learned come from research and from building. We would flag clearly that none of them is validated by a buyer:

1. The market is barbelled, and the middle is structurally empty. Enterprise platforms start around \$100k/yr; below that there is nothing but Google Alerts. Incumbents cannot descend without breaking their cost structure, so the gap persists rather than closing on its own. That is our entire thesis — and the thing we most need customers to confirm or destroy.

2. Trajectory matters more than the score. A supplier rated 61 tells you little. A supplier that went 82 → 74 → 61 over three weeks tells you to call them today. This shaped the product decisively: daily rescanning and stored history, rather than one-off assessments.

11. What proof do you have that customers will actually pay for this — not just say they like it?

No direct proof. No paying customers, no letters of intent, no pilots. We will not dress up inference as evidence. Obtaining that proof is the single most important goal of our next six months, and the main reason we are applying to a programme rather than only raising capital.

What we have is indirect evidence that the job is worth paying for:

Willingness to pay is proven at the top of the market. Resilinc, Interos, Everstream and Prewave sustain six-figure annual contracts doing this job for large enterprises. The value of the outcome — knowing a supplier is failing before the shipment doesn't arrive — is well established. What is unproven is whether a 100-person manufacturer will pay a small monthly fee for a lighter version. That is precisely the open question.

The cost of the alternative is orders of magnitude higher. A single disruption — expedited air freight, a spot-market buy, a stopped production line — routinely costs tens of thousands of dollars. Against that, a few hundred dollars a year is a rounding error. Favourable economics do not guarantee a sale, but they make price unlikely to be the objection.

Regulation is creating a non-discretionary reason to buy. Germany's LkSG is live and the EU's CSDDD is phasing in, pushing documented supplier due diligence onto companies well below enterprise size — buyers who currently have no affordable way to comply. This reaches Sri Lankan exporters directly, as Western customers increasingly ask them to evidence their own supply chains.

What we consider real proof, and how we will get it: not survey responses, not "this looks useful", and not letters of intent. Payment. Our target is at least five paying customers at validated pricing within six months, drawn from ten design partners.

And if it does not happen: if mid-sized buyers will not pay after twenty serious conversations, that is a genuine signal about the business, and we would rather learn it in month four than month twenty. Our low cost base means we can afford to test the question honestly rather than needing the answer to be yes.

C. Your product

What you have built and what makes it different.

12. Describe your product in one simple sentence anyone could understand.

A credit score for a business's suppliers — updated automatically by AI, with alerts before problems hit the business.

13. Where is your product today — an idea, a prototype, or already used by customers? You can paste a link and/or attach images below.

A working prototype, feature-complete and running end-to-end, but not yet used by customers.

Functioning today: enter a supplier name and the AI agent runs 5 parallel web searches, extracts risk signals, synthesises a 0–100 health score with supporting evidence, and returns it in under 60 seconds. A scheduled job rescans every supplier automatically, building the score history behind the trajectory charts. Email and Slack alerts fire when a score crosses a customer-set threshold. Authentication, the full dashboard, supplier management, the signals feed, and Stripe billing across three tiers are all built and working.

What's missing is not product — it's customers. We have not launched publicly, run a pilot, or taken a payment. The technical risk is largely retired; the entire remaining risk is commercial.

14. What is the key idea that makes your product different from others?

Incumbents own the risk data. We don't need to.

Enterprise platforms license news feeds, buy credit databases, map your supply chain and onboard your suppliers. That data pipeline is the six-figure price tag and the six-week implementation. An AI agent reading the open web on demand removes it entirely.

Three consequences, and they are the business:

Zero onboarding. A supplier name is the entire setup — score in 60 seconds, not six weeks. This removes the single biggest conversion barrier for a smaller buyer.

A price point nobody else can reach. No data licensing means a low monthly fee is viable, which is what makes the un-served mid-market reachable at all. Our Sri Lankan cost base compounds this: we can profitably sustain a price that a US or European competitor would lose money at.

Coverage of small private suppliers with no credit file or database entry — usually the risky ones, and precisely where data-licensing approaches go blind.

The second idea, which drives retention: trajectory, not snapshot. We rescan daily and keep the history, so customers see 82 → 74 → 61 rather than a single number. That history cannot be backfilled by a competitor a customer switches to, so it compounds into a switching cost. Every score also retains its underlying signals and sources — a procurement lead will not act on a number they cannot defend to their boss.

15. How is it built or delivered, and what outside tools or suppliers do you rely on?

A thin serverless application on rented infrastructure — no servers, no data centre, no ops team. A Next.js web app plus around 15 serverless functions on Netlify, with Supabase for authentication and Postgres.

The core is a single pipeline run per supplier: idempotency check → 5 parallel web searches → AI signal extraction → score synthesis → persist → alert if the threshold is crossed. A nightly cron job scans every active supplier.

External dependencies:

OpenAI — GPT-4o-mini extracts signals, GPT-4o synthesises the score.

Tavily, with Serper as automatic fallback — web search.

Supabase — authentication and Postgres.

Netlify — hosting, serverless functions, scheduled jobs.

Stripe — billing. Resend and Slack webhooks — alert delivery.

Two points worth noting. No vendor is a moat, and no vendor is a hostage — every dependency is commodity and substitutable, and search is already dual-sourced with live failover, since an outage at a single search provider would otherwise stop the product. Second, there is no data licensing contract anywhere on that list. That is the entire cost-structure argument: incumbents are locked into feed licenses and implementation staff, while our cost is metered API calls that scale to zero.

Cost per scan is a few cents — the cheap model does the bulk work and the expensive one only the final synthesis. Per-run AI spend is logged in the code, so margin is measured rather than estimated.

16. What are you deliberately NOT building in the next 18 months?

1. N-tier supply chain mapping — knowing your supplier's suppliers. The marquee enterprise feature, and it requires exactly the data pipeline we refuse to build. Chasing it makes us a worse version of Resilinc.

2. ERP and procurement integrations (SAP, Coupa, NetSuite). Every enterprise prospect asks; each is months of work and pulls us into nine-month sales cycles with a buyer we are not built for. Our customer's system of record is a spreadsheet.

3. Enterprise readiness — SSO, SOC 2, RBAC, custom contracts. All of it exists to close deals we have chosen not to chase yet. We will build it when a specific customer is paying us to.

4. Buying proprietary data. Licensing a credit or news feed would make scores look more authoritative and would reintroduce the fixed cost that prices incumbents out of this segment. Zero data licensing is the strategy, not a gap.

5. Training our own model. We are a workflow and data-history company, not an ML company. Frontier models get cheaper and better each quarter at no cost to us.

6. Adjacent risk categories — ESG scoring, compliance reporting, forced-labour audit trails. Real markets with genuine regulatory tailwind, but a different buyer and a different sale. Taking them on before the core product retains would fragment a three-person team.

Instead: make daily scanning accurate enough and alerts trusted enough that a procurement lead checks the dashboard on Monday morning — then prove retention before adding anything.

The trigger that would change this: if inbound demand from 500+ employee companies becomes the majority of signups, items 2 and 3 move up. We would rather be pulled up-market by revenue than push there on a hunch.

D. Your market

How big is the opportunity?

17. How big is your market? Give your best estimate and explain how you worked it out.

Best estimate: ~\$115M TAM at today's pricing, ~\$285M once repriced, with a credible path to \$1–2B.

Bottom-up. Our buyer is a manufacturer or distributor with 20–500 employees in the US, UK, EU, Canada and ANZ:

US manufacturers, 20–500 employees: ~60,000

US wholesale distributors: ~40,000

UK + EU equivalent: ~150,000 (manufacturing-heavy in Germany and Italy)

Canada + ANZ: ~20,000

Total qualifying firms: ~270,000

Applying revenue per customer. At current pricing (Pro \$19.99/mo, Enterprise \$99.99/mo) blended ACV is ~\$430, giving ~\$115M. Our unit-cost analysis shows that pricing sits below cost to serve, so we plan to reprice to ~\$49 / \$249. That lifts blended ACV to ~\$1,050 and TAM to ~\$285M.

Cross-check, top-down. Supply-chain risk software is a ~\$3–4B category. A ~\$285M under-served SMB tier inside it is roughly 8% — consistent, so the two methods agree.

SAM (reachable self-serve, no sales team — ~40% of qualifying firms): ~\$115M.

SOM, 3 years: 1,500–3,000 paying customers → \$1.5–3M ARR. That is the number to hold us to.

Why this is a large opportunity for a Sri Lankan company specifically. \$1.5–3M ARR would be a modest outcome for a Silicon Valley startup burning \$300k a month. At our cost base it is a genuinely profitable, self-sustaining software business earning in hard currency — and it is reached with roughly a tenth of the capital a US competitor would require.

We would rather flag the limitation than inflate it. \$285M is modest for a venture-scale outcome. Expansion comes from two levers, not customer count: pricing (per-supplier and seat pricing moves blended ACV toward \$3–5k, taking the same base past \$800M), and widening beyond manufacturing to any business with suppliers — healthcare, construction, retail — a further 3–5x with no product change. Both together reach \$1–2B. Neither is a claim; they are what the next 18 months are designed to test.

Where these numbers are soft: firm counts are order-of-magnitude estimates from public statistics we would want verified against Census and Eurostat data. The weakest assumption is ACV — a price we set, not one the market has confirmed.

18. Which countries or regions will you focus on first, and why those?

Sri Lanka as the proving ground; the United States and United Kingdom as the scale market.

Sri Lanka first, for validation rather than revenue. Our design partners should be within reach, and Sri Lanka's export manufacturers — apparel, tea, rubber, food processing — have exactly this problem twice over. They depend on upstream suppliers of fabric, chemicals and packaging, and they simultaneously sell to Western brands who are themselves being pushed by EU regulation to document supplier due diligence. We can meet these companies face to face, iterate weekly, and learn to sell in our own market before selling remotely. Acorn Group's own supplier relationships would be an ideal starting point.

United States as the primary revenue market. The largest concentration of mid-sized manufacturers and distributors, the highest willingness to pay for SaaS, English-language news sources our agent reads natively, and self-serve card-based purchasing that suits a product-led motion with no sales team. Software has no geographic delivery constraint — we can serve US customers from Colombo with no disadvantage.

United Kingdom second — same language, same buying behaviour, small enough to test cheaply, and a favourable time-zone overlap with Sri Lanka for support.

Germany third, deliberately timed. The densest population of qualifying Mittelstand manufacturers in Europe, with a hard regulatory forcing function in the Supply Chain Act (LkSG). We are holding it until we can support German-language sources — entering under-resourced would waste the strongest tailwind available to us.

India and South Asia are a live option we are not yet committing to. Enormous manufacturing base and regional proximity, but lower willingness to pay. We will test it only if the US self-serve motion proves harder than expected.

19. Is this market growing, flat, or shrinking? What is driving that?

Growing — high single to low double-digit annual growth in the underlying software category, with the SMB tier growing faster because it starts from near-zero penetration.

Structural. Supply chains became longer, more global and more single-sourced over two decades, permanently raising the cost of any one supplier failing.

Event-driven. The pandemic, Suez, the Red Sea, tariff volatility and sanctions turned supply disruption from a tail risk into a recurring one. Budget for this category is now defended rather than questioned.

Regulatory — the driver that most directly creates our segment. Germany's LkSG is live and the EU's CSDDD is phasing in, pushing documented supplier due diligence onto companies well below enterprise size. This runs on a legislative timetable, not a news cycle. It also reaches Sri Lanka directly: local exporters selling into Europe are increasingly asked to evidence their own supply chains, which gives us a regionally relevant wedge as well as a Western one.

The honest counter-argument: some post-2020 urgency fades as memories do. Our position does not depend on panic — the regulatory driver is structural and dated.

20. Are there any licences, rules, or regulations you need to follow?

No licence is required to operate. We are not regulated as a credit rating agency or financial adviser, and we take care to stay outside those definitions: our output is an operational indicator for a buyer's own sourcing decisions, not a creditworthiness rating sold to investors. Every score carries a disclaimer that it is an indicator only — not financial, legal or investment advice.

Obligations arising from our customers' jurisdictions:

Data protection (GDPR / UK GDPR). We hold customer account data and process supplier information. We operate under data-processing agreements with our sub-processors (Supabase, OpenAI, Stripe). As a non-EU company serving EU customers we will likely need an EU representative and appropriate transfer safeguards — an area where we would value guidance.

Defamation and accuracy — our most material legal exposure. We publish negative assessments of named companies. We mitigate by citing sources for every signal, storing the evidence behind each score, keeping outputs factual rather than characterising, and offering suppliers a correction route.

EU AI Act. Our use sits outside the high-risk categories, but transparency obligations apply — users must know they are seeing AI-generated output. We disclose this already.

Payments. Stripe handles all card data, keeping PCI scope out of our systems.

Sri Lankan obligations:

Personal Data Protection Act No. 9 of 2022 applies to our handling of personal data as a locally established company, and we will need to align our practices with it as it comes fully into force.

Foreign exchange and inward remittance. Revenue will arrive in USD from overseas customers. We will need appropriate banking arrangements and to comply with Central Bank rules on inward remittances and foreign-currency accounts, and to confirm whether Stripe can settle to a Sri Lankan entity or whether an offshore structure is required. This is a practical question we would welcome Acornic's help resolving — it affects entity structure and is not something we can answer ourselves.

Tax treatment of export services. Software delivered to overseas customers should qualify for export-service treatment, but we will confirm the current position with local counsel.

Not legally required but commercially necessary as we move up-market: SOC 2 Type II. Deliberately deferred until a paying customer requires it.

E. How you make money

Your pricing and the economics.

21. How do you make money? What do you charge, and how?

Monthly SaaS subscription, self-serve, charged by card through Stripe in USD. Three tiers gated on supplier count and scan frequency:

Starter — free. 3 suppliers, weekly scans, email alerts. An acquisition channel, not a revenue line.

Pro — \$19.99/mo. 25 suppliers, daily scans, email + Slack alerts, score trajectory, signals feed, on-demand scans. The core product.

Enterprise — \$99.99/mo. Unlimited suppliers, 6-hourly scans, priority support.

No setup fees, no professional services, no per-seat charge. Revenue grows two ways: free-to-Pro conversion, and Pro-to-Enterprise upgrade as a customer's supplier list outgrows 25.

The pricing metric is supplier count — the right choice, since it tracks both the value the customer receives and our cost to serve.

All revenue is earned in USD from overseas customers, while our costs are largely in LKR. That currency asymmetry works in our favour and gives us margin resilience a US competitor does not have.

We expect to reprice to ~\$49 / \$249 before launch. Our own unit-cost analysis (Q22) shows current pricing sits below cost to serve. Validating the right price with design partners is a first-quarter priority.

22. For a typical customer, what does it cost to win them, what do they pay you, and what is your profit margin? If you're still pre-revenue, your best estimate is fine.

Pre-revenue, so these are estimates with the working shown.

Cost to win (CAC): product-led and self-serve — SEO, content and founder-led outbound, no sales team. Estimated \$150–350 blended per paying customer, assuming free-tier signups convert at 3–5%.

What they pay: \$240/yr on Pro today, \$588/yr repriced. At an assumed 24-month retention, gross LTV is ~\$480 today, ~\$1,175 repriced.

Margin — a real problem we have identified rather than one we are hiding. Each supplier scan costs roughly \$0.06: five web searches (\$0.03), signal extraction on the cheap model (\$0.003), score synthesis on the expensive model (~\$0.025). Per-run spend is instrumented in our code, so this is measured, not guessed. Note these are USD API costs — the one part of our cost base that gets no benefit from being in Sri Lanka.

A Pro customer watching 25 suppliers daily generates 750 scans/month — about \$45 of cost against \$19.99 of revenue. Even a realistic 10-supplier customer costs ~\$18/mo to serve. As currently priced, Pro is gross-margin negative.

Three fixes, in order of impact:

1. Deduplicate scans across customers. A supplier's risk score is not customer-specific. Today each customer's supplier is scanned separately; scoring canonical companies once and fanning the result out means cost scales with distinct suppliers, not customers. At any density that is a 5–20x cost reduction — and margin then improves as we grow.

2. Reprice to ~\$49 Pro / ~\$249 Enterprise. Defensible against a six-figure enterprise alternative.

3. Scan smarter — cheaper models for extraction, event-triggered rather than blind daily rescans.

With those in place: contribution margin of 75–85%, LTV/CAC of 4–6x, CAC payback in 4–7 months. Without them the model does not work — which is why it is the first thing the next six months must prove.

23. How and when will the business become profitable?

This is where being a Sri Lankan company changes the answer fundamentally.

Our fixed costs are three founder salaries at local rates, plus infrastructure. No offices, no servers, no sales team. Fully operational with all three founders full-time, our total burn is approximately \$4,500–5,500 per month, rising to about \$7,000 once we add a part-time commercial hire. An equivalent US-based team would burn five to six times that.

At repriced levels and post-deduplication cost, each Pro customer contributes roughly \$35–40/month. Covering a \$7,000 monthly cost base therefore requires approximately 175–200 paying customers — not the 800-plus a Silicon Valley team would need.

We expect operating breakeven 18–24 months from launch, gated by acquisition rate rather than by cost.

The strategic consequence: we can become default-alive on a single small round. That is not a rhetorical point — it means the next raise becomes a choice about growth speed rather than a matter of survival, and it means a modest ARR outcome is a genuinely profitable business for us where it would be a failure for a US-funded competitor.

Two things move the date: free-to-paid conversion rate, and whether scan deduplication ships before free-tier volume becomes expensive.

24. Do customers pay upfront or later? Do you need to hold stock or pay for things well before you get paid?

Customers pay upfront, monthly, by card via Stripe. No invoicing, no net-30 terms, no receivables risk, no collections. We plan to offer annual prepayment at a discount, which improves cash further.

No inventory, no stock, no capital equipment, nothing physical to finance.

Working capital runs negative in our favour. We collect at the start of each month while our vendors — OpenAI, Tavily, Netlify, Supabase — bill monthly in arrears. Cash arrives before the cost of serving it is paid.

The one genuine cash drag is the free tier. Free users cost real money: 3 suppliers scanned weekly is roughly \$0.78/month per free account in API spend. At 10,000 free accounts that is ~\$7,800/month before any of them convert. It is the main reason scan deduplication matters, and why free-tier scanning runs weekly rather than daily.

One operational point specific to our situation: revenue arrives in USD via Stripe while most costs are in LKR. We will need banking arrangements that allow us to receive and hold foreign currency efficiently, and to confirm Stripe settlement to a Sri Lankan entity. We have flagged this in Q20 as an area where we would value guidance.

F. Progress so far

What you have achieved to date.

25. If you have early numbers (users, sign-ups, paying customers, revenue, repeat usage), share them and the trend. If you're pre-launch, tell us what you've built or tested so far.

No revenue or users yet — we're pre-launch. What exists is a working end-to-end product, not a mockup or a deck.

Built and functioning today:

The scoring agent. Enter a supplier name and a 6-step pipeline runs: idempotency check, 5 parallel web searches, AI signal extraction, score synthesis, persistence, and alerting. Returns a 0–100 health score with supporting signals in under 60 seconds.

Daily automated monitoring. A scheduled job rescans every active supplier on a cron, building the score history that trajectory charts are drawn from.

Threshold alerting. Email (Resend) and per-supplier Slack webhooks fire when a score drops past a customer-set threshold.

Full product surface. Supabase authentication, supplier CRUD, dashboard, trajectory charts, signals feed, alerts history, and live scan-progress polling.

Billing. Stripe checkout, billing portal, and webhooks are wired end-to-end, with the three tiers enforced in code.

Production hardening. All external web content is sanitised before entering model prompts (prompt-injection defence), search is dual-sourced with automatic failover, and per-run AI spend is logged so unit economics are measured rather than guessed.

26. Do you have any early customers, users, or partners? If so, who — and can we speak with them? (It's fine if you don't have any yet.)

Not yet. We've deliberately built to a working product first so that early conversations can start with a live score for a prospect's own supplier rather than a description of one.

Recruiting 5–10 design partners is the immediate next step, targeting mid-sized manufacturers and distributors. Happy to share that list as it forms.

27. Do you have any signed deals, pilots, or letters of intent? (you can attach them separately)

None. Nothing has been offered to a customer yet, so there is nothing signed or in negotiation.

28. What have you tried that did not work, and what did you learn?

Nothing has been tested in the market yet, so these are engineering and economic learnings from the build — the market learnings are still ahead of us.

Running everything on the top model was unaffordable. GPT-4o across the whole pipeline made per-scan cost impossible at SMB pricing. We split the work: the cheap model does the bulk extraction, the expensive one only the final synthesis. Roughly 5x cost reduction with no quality loss. Lesson: model selection per step is a pricing decision, not an engineering detail.

A single search provider wasn't reliable enough. Outages and rate limits took scans down entirely. We added a second provider with automatic failover. Lesson: at the point where an external API is your product, one supplier is a single point of failure.

Scoring was too slow for a normal request/response. A 60-second pipeline times out or looks broken behind a spinner. We moved to asynchronous scanning with client-side progress polling. Lesson: the perceived speed of an AI product is a UX problem as much as a latency one.

Feeding raw web content into prompts was a security hole. Scraped pages can contain instructions that hijack the model. Everything external now passes through sanitisation and is wrapped in delimiters before reaching a prompt. Lesson: in an agent that reads the open internet, every input is untrusted.

Fanning out daily scans naively hit rate limits. Scanning all suppliers at once got us throttled. Now batched with pacing between groups.

The biggest one — our launch pricing doesn't cover our costs. Instrumenting real per-scan spend showed that Pro at \$19.99 with daily scanning of 25 suppliers is gross-margin negative. Two fixes follow directly: reprice, and deduplicate scans so a supplier watched by many customers is scored once rather than once per customer. Lesson: instrument unit costs before setting prices, not after — and in AI products, the architecture determines the margin.

G. Competition

Who else does this, and why you will win.

29. Who are your main competitors (3–5)? Where do they beat you, and where do you beat them?

1. Resilinc / Interos / Everstream Analytics — enterprise supply-chain risk platforms, \$100k–500k/yr contracts. They beat us on: depth — n-tier mapping, licensed proprietary data, human analyst validation, enterprise trust and certification. We beat them on: price and time-to-value. They need weeks of onboarding and a dedicated analyst; we need a supplier name and 60 seconds. They structurally cannot serve a 25-supplier company — their cost base is data licensing and implementation staff.

2. Prewave — Vienna-based, well-funded, strong EU regulatory positioning. They beat us on: multilingual coverage at scale, established LkSG/CSDDD compliance credibility, real enterprise customers, funding. We beat them on: self-serve. They sell top-down to large European manufacturers; a 60-person distributor has no way to sign up and get a score today.

3. Craft.co — supplier intelligence and risk data with mid-market reach. They beat us on: breadth of company data — financials, headcount, ownership. We beat them on: focus. They are a data platform you must interpret; we deliver one score, a trend, and an alert that tells you when to act.

4. Dun & Bradstreet / Moody's supplier risk modules. They beat us on: brand trust and credit data nobody else has. We beat them on: coverage of the small private suppliers with no credit file — usually the risky ones — and on the operational signals (fires, strikes, recalls) that credit data misses entirely.

5. The real competitor: doing nothing, or Google Alerts and a spreadsheet. It beats us on: being free and frictionless. We beat it on: actually working — alerts have no score, no history, no threshold, and no signal-to-noise. Most of our sales effort goes here. In practice we compete with inertia, not with Resilinc.

On being Sri Lankan: every named competitor is US or European with a matching cost base. None has a local presence in South Asia, and none is defending the price point we intend to occupy. We do not expect our location to be a disadvantage in a self-serve software market — and it is a decisive advantage on cost.

30. What will stop others from copying you in 2 years? And in 5 years?

Frankly: With a 2-year-old system, not very much technically. The pipeline itself consists of search plus two model calls. A good team can remake the core in a couple of months. We prefer to say that than pretend we have a moat that we don't.

What protects us over 2 years is position and workflow, not code: the accumulated score history for each customer's suppliers — which begins the day they sign up and cannot be backfilled by a competitor they switch to — the alert routine embedded in a procurement lead's week, and being the first recognizable name in a segment nobody currently serves.

At 5 years, three things compound into a real moat:

Longitudinal data. Score histories paired with actual outcomes — which suppliers really failed — let us validate and tune predictive accuracy. That dataset only accrues with time. A new entrant in 2031 starts with zero history and no way to prove their scores mean anything.

Cost density. A supplier's risk profile is not customer-specific. Once scanning is deduplicated, the marginal cost of serving the hundredth customer watching a common supplier approaches zero. Coverage density becomes a structural cost advantage — and stacked on top of a Sri Lankan operating base, it lets us hold a price point a competitor loses money at.

Distribution and workflow depth. Alert routing, Slack integration and eventually procurement-system integrations make us the tool the team already uses. Switching means losing history and rebuilding a routine.

31. If a big, well-funded company tried to copy you, what would they need to do — and how long would it take?

Technically: 3–6 months. The pipeline is not the hard part. A well-resourced team could ship a functional clone within two quarters, and would likely engineer it better than we have.

What they would need beyond the build: a self-serve motion (most incumbents have none — they sell through enterprise reps), SMB-grade pricing and support economics, a brand SMBs will buy from, and the willingness to run a \$49/mo product alongside a \$200k one.

That last point is the real barrier, and it is not about time. An incumbent launching a cheap self-serve product creates immediate channel conflict: their own sales team competes with it, and existing customers ask why they pay 200x for something adjacent. Their cost structure — data licences, implementation staff, enterprise sales — makes SMB pricing structurally unprofitable. This is textbook disruption: they can see it clearly and still rationally decline to act.

For big tech (SAP, Microsoft, Google), the market is simply too small to warrant attention until it is not — at which point it is an acquisition target rather than a build.

Where we are genuinely exposed: another startup copying us. No channel conflict, no legacy cost base. Our defence there is not technology — it is moving faster, getting the score-history flywheel turning first, and owning the mid-market position before someone else claims it. Our low burn helps here in a specific way: we can survive a long enough search for product-market fit that a US-funded competitor with a 12-month runway cannot.

H. Money & funding

Your finances and what you are raising.

32. If you're already trading, summarise your last 12 months — revenue, cash in the bank, monthly spend, and runway. If you're pre-revenue, just tell us your current cash position and monthly spend.

Pre-revenue. No trading history, no revenue, no external capital raised. Entirely founder-funded.

Cash position: approximately [USD X / LKR X], contributed by the three founders. No debt or liabilities.

Monthly spend: approximately [USD 50–150]/month — AI and search API usage during development (OpenAI, Tavily), hosting (Netlify, Supabase), domain and tooling. All three founders are unpaid and part-time, so there is no salary cost in that figure.

Runway: effectively open-ended at current burn, since we have no fixed obligations — but that is also the constraint. At this level we can keep the product running; we cannot go full-time or acquire customers. This raise converts a working product into a commercial one.

Capital efficiency to date: we built a complete working system — AI scoring agent, automated daily monitoring, alerting, full web application, Stripe billing — for under [USD X] in total. That is the strongest evidence we can offer about what we do with money.

33. What's your rough plan for the next 12–24 months — revenue, costs, team size, and cash? Best estimates are fine at this stage.
Assumes repricing to ~\$49 / \$249, which the unit economics require, and all three founders full-time on investment.

Months 0–6 — validate. Recruit 10 design partners, starting in Sri Lanka where we can meet buyers face to face, ideally including companies in the Acorn Group ecosystem. Confirm pricing against real willingness to pay. Ship scan deduplication to fix gross margin. Convert first paying customers. Target: ~15 paying customers, ~\$700 MRR. Team: 3 founders. Burn: ~\$4,500/mo.

Months 6–12 — find the acquisition channel. Turn to the US and UK self-serve motion. Add a part-time growth hire. Build SEO and content around supplier risk; test outbound to procurement leads. Target: ~100 paying customers, ~\$4,500 MRR by month 12. Burn: ~\$6,500/mo.

Months 12–18 — scale what works. Double down on the one channel that converts. Consider German-language coverage if the EU regulatory pull proves real. Target: ~250 customers, ~\$11,000 MRR. Team: 3 founders + 1 growth. Burn: ~\$7,500/mo.

Months 18–24 — default alive. Target: ~450 customers, \$20,000 MRR (\$240k ARR), contribution margin 75–85%, team of 4–5. Operating cash-flow breakeven around month 20–22 — earlier than a comparable Western team, because our cost base is a fraction of theirs.

Cash: the raise plus growing revenue covers the full period with roughly 4–5 months of buffer. The plan is designed to reach breakeven on this round, not to require a bridge.

The assumption most likely to be wrong: free-to-paid conversion, assumed at 3–5%. If it comes in at 1%, customer targets halve and we would cut spend rather than chase volume. Our low burn means that adjustment is survivable rather than fatal.

34. How much are you raising from Acornic, in what form (equity / SAFE / convertible), and what will you use it for?

Raising USD 100,000 as a pre-seed investment. We are open to whichever instrument Acornic prefers — SAFE, convertible note, or priced equity. Our preference is a SAFE with a post-money valuation cap, since it keeps the round fast and inexpensive and avoids fixing a valuation before we have revenue data, but we will follow Acornic's standard terms.

The round funds 18 months to commercial validation. Because our cost base is Sri Lankan, this goes considerably further than the same amount would for a US or European team — a structural advantage we intend to press.

Item Amount Purpose

Founder salaries — 3 × 18 months ~\$43,000 All three full-time; currently unpaid and part-time

Marketing & customer acquisition ~\$22,000 SEO, content, paid experiments to find one repeatable channel

AI, search & infrastructure ~\$15,000 Scan costs including free-tier usage as volume grows

First commercial hire (part-time, from month 9) ~\$10,000 Growth/sales — the capability we lack

Legal, incorporation, accounting & tools ~\$10,000 Entity setup, IP assignment, cross-border structuring

Total ~\$100,000 18 months

What it buys: roughly 175–200 paying customers and operating breakeven within 18–24 months, making the next raise a choice about growth speed rather than survival.

What it is not for: building more product. The scoring agent, automated monitoring, alerting and billing all work today. This capital buys time and distribution to answer the only question that still matters — whether the market wants it.

Why Acornic specifically. We are three developers with no commercial experience. Acornic's four-week sprint, buyer showcases and mentoring address our actual weakness rather than one we do not have. And Acorn Group's own operations depend on suppliers — we would welcome the group as a first design partner, giving us honest feedback and a reference customer in a single step.

35. Are any other investors already involved or in discussion?

No. The company is entirely founder-funded, with no outside investors, no advisers holding equity, no grants, and no other conversations in progress. Acornic is the first investor we have approached — we sought a programme with hands-on commercial support and local presence rather than capital alone, and Acornic is the clearest fit in Sri Lanka.

36. What specific goals will this funding help you reach?

Five measurable milestones, each retiring a specific risk:

1. Prove people pay — 10 design partners converted to paying customers within 6 months. Retires the core risk: that supplier monitoring is a nice-to-have below enterprise size. We would start with Sri Lankan export manufacturers, where we can meet buyers in person and iterate weekly.
2. Fix gross margin — ship scan deduplication and validated pricing, reaching 75%+ contribution margin. Our largest known weakness. Until this is fixed, growth makes losses larger rather than smaller.
3. Find one repeatable acquisition channel — CAC under \$350 with 3%+ free-to-paid conversion, by month 12. The only unresolved question that funding directly buys an answer to.
4. Reach ~\$20k MRR / ~\$240k ARR by month 24, with annual retention above 85%.
5. Reach operating breakeven by month 20–22, so the next round is optional. At our cost base this requires roughly 175–200 customers, which we consider genuinely achievable on a single \$100k round.

What this funding is explicitly not for: building more product. The technical risk is largely retired — the agent, monitoring, alerting and billing all work today. This capital buys the commercial capability we do not have, and the time to use it.

I. Fit with the Acornic programme

Acornic backs early-stage founders with an accelerator / incubator programme — tell us how we can help.

37. Beyond funding, what support would help you most from Acornic's programme? (e.g. mentoring, product, sales & marketing, hiring, introductions, regulatory help)

In priority order, shaped by the honest gap in our team — we are three developers with no commercial background.

1. Sales and go-to-market — by far the most valuable thing Acornic can give us. We have a working product and no proven way to reach a procurement manager at a 100-person manufacturer. We need help designing the acquisition motion: outbound that works in an unglamorous B2B segment, pricing and packaging validation, and how to run a design-partner programme properly. Acornic's four-week sprint model is aimed squarely at this, and it is exactly where we are weakest.
2. Customer introductions — and specifically, Acorn Group itself. Acorn Group has real operations that depend on suppliers. We would value the group as a first design partner: it would give us honest feedback from a genuine buyer and a reference customer in the same step. Beyond that, introductions to manufacturers, distributors, and procurement leaders in the Acornic network. Ten real conversations would teach us more than another quarter of building.
3. Your industry-buyer showcases. Access to curated buyer audiences is worth more to us than demo-day exposure to investors. Our constraint is buyers, not capital.
4. Mentoring from someone who has sold B2B SaaS internationally. Our market is the US and UK while our team sits in Colombo. We need guidance on selling remotely and cross-border: conversion benchmarks, trial design, and how to price against a six-figure enterprise alternative.
5. Commercial hiring and regulatory guidance. Our first non-founder hire should be growth, not engineering, and we don't yet know how to evaluate that person. We would also value help on entity structure for serving US/EU customers from Sri Lanka, and on positioning against EU supply-chain regulation without overclaiming compliance coverage.

38. Are the founders full-time on this — or ready to be? Can you fully commit to taking part in the programme?

All three founders currently work part-time on SupplyPulse alongside current employment & studies. All three are ready and committed to going full-time on investment, and that transition is built into our use of funds.

We can fully commit to the programme and will treat it as our primary focus throughout. Being Colombo-based, we can participate in person.

39. Is the company already set up as a legal entity? If so, where and when? If not, what's your plan?

Not yet incorporated. We deferred it to avoid ongoing compliance costs while pre-revenue and unfunded.

Our plan: incorporate a private limited company in Sri Lanka at or before the point of investment, with the three founders as equal shareholders, standard four-year vesting with a one-year cliff, and full IP assignment from each founder to the company.

One open question we'd welcome guidance on: our customers will be primarily in the US and UK, paying by card in USD. We would value Acornic's view on whether a Sri Lankan entity alone is the right structure, or whether an offshore holding company is worth establishing early for payment processing, customer contracting and future fundraising. We'd rather get this right at the start than restructure later.

40. Why are you the right team to build this, and why now?

Why us. Three developers who have already built and shipped the thing. This is not a concept — the AI scoring agent, automated daily monitoring, alerting, the full application and Stripe billing are live and working. We built all of it unfunded and part-time, which says something about what we can do per dollar.

The product is fundamentally an engineering problem in commercial clothing: an AI agent that reads the open web reliably, cheaply and safely at scale. Making that work — prompt-injection defence against untrusted web content, dual-sourced search with automatic failover, cost-optimised model routing, instrumented unit economics — is precisely what three engineers are good at. Most teams attacking this segment would have to hire that capability. We are it.

Our Sri Lankan cost base is a structural advantage, not a limitation. Our competitors serving this market are US and European companies with US and European burn rates. We can reach profitability on a fraction of the capital and can sustainably serve a price point they cannot. Software sold to US and UK buyers has no geographic constraint on delivery.

Our honest weakness, stated up front: none of us has sold B2B software. That is a real gap, and it is exactly why we are applying to a programme rather than only raising money.

Why now. Three things converged. Frontier AI model costs fell far enough that scoring a supplier costs cents rather than dollars — this business was not viable at 2023 prices. Supply-chain disruption became a recurring operational reality rather than a tail risk, so the budget conversation has changed. And European regulation — Germany's LkSG now live, the EU's CSDDD phasing in — is pushing documented supplier due diligence onto companies far below enterprise size, creating our exact segment on a legislative timetable rather than a hopeful one.

The window is open now. Someone will serve the mid-market. We would like it to be us.

41. What would you most want to achieve in the next 3–6 months with Acornic's help?

One primary goal: prove that mid-sized companies will pay for this. Ten design partners using the product weekly, with at least five converting to paid at validated pricing. Nothing else matters until this is answered — and it is a commercial question we currently lack the skills to answer alone. That is the core of what we want from the programme.

Supporting goals:

Get out of the building, immediately. Twenty structured discovery conversations with procurement and operations leaders — ideally starting inside the Acorn Group ecosystem — within the first three months. We have spoken to zero customers to date. Correcting that is urgent.

Fix the unit economics. Ship scan deduplication and reprice, taking contribution margin from roughly break-even to 75%+. We know exactly what to build; we need help knowing what to charge.

Find the first acquisition channel. Run three or four structured experiments — SEO, outbound, partnerships, communities — and identify one with customer acquisition cost under \$350. Not scale it; just find it.

Learn to sell. By month six, at least one founder should be able to run a discovery call and close a deal unaided. In a three-developer team, that single capability is the difference between a good product and a company.

What success looks like at month six: first paying customers, positive contribution margin, one acquisition channel showing early signal, and three founders materially less naive about selling than we are today.